



CS5250 ASSIGNMENT 4

YANG MO – A0091836X

Task 1

1. Screenshot of the github commit & Code

Please see **appendix A** for github screenshot and **appendix B** for device driver code

Task 2

1. New kmalloc and kfree code in init and exit function

- kmalloc:

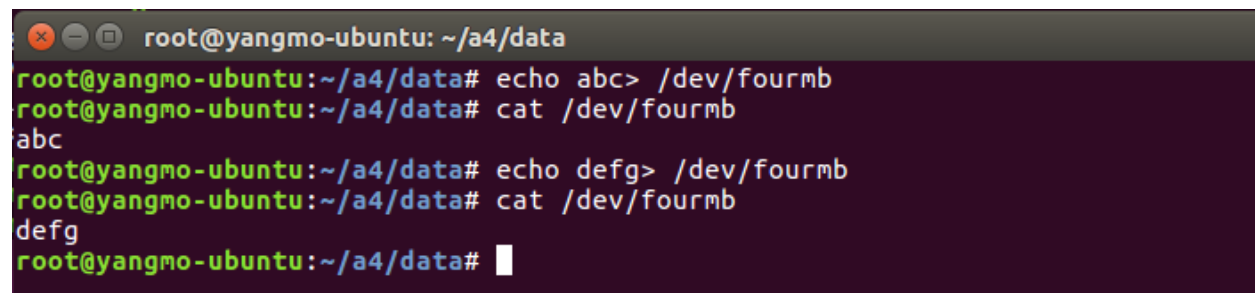
```
// allocate 4 MB of memory for storage
// kmalloc is just like malloc, the second parameter is
// the type of memory to be allocated.
// To release the memory allocated by kmalloc, use kfree.
fourMB_data = kmalloc(1024*1024*4*sizeof(char), GFP_KERNEL);
```

- kfree:

```
// if the pointer is pointing to something
if (fourMB_data) {
    // free the memory and assign the pointer to NULL
    kfree(fourMB_data);
    fourMB_data = NULL;
}
```

2. Output of the test cases:

As can be seen below, new device can take more than 1 byte:

A terminal window titled 'root@yangmo-ubuntu: ~/a4/data' showing two successful writes to a device. The first write is 'abc' and the second is 'defg'.

```
root@yangmo-ubuntu: ~/a4/data
root@yangmo-ubuntu:~/a4/data# echo abc> /dev/fourmb
root@yangmo-ubuntu:~/a4/data# cat /dev/fourmb
abc
root@yangmo-ubuntu:~/a4/data# echo defg> /dev/fourmb
root@yangmo-ubuntu:~/a4/data# cat /dev/fourmb
defg
root@yangmo-ubuntu:~/a4/data#
```

3. Copy a 5MB file to your device and display total bytes written and show unwritten part I am using a printk function in write function to print the total bytes written so far:

```
printk("Total Bytes written so far: %d ", bytes_written_total);
```

This `bytes_written_total` will record the latest total written bytes so far since the open function is called. It will be reset to zero each time open function is called and will accumulate if write function is called multiple times to write an large file (such as in our case 4MB or 5MB file).

- Create exactly 4MB file:

Create exactly 5MB file:

Created files:

```
root@yangmo-ubuntu: ~/a4/data
root@yangmo-ubuntu:~/a4/data# ls -lh
total 9.0M
-rw-r--r-- 1 root root 5.0M Apr  2 21:56 fivemb.txt
-rw-r--r-- 1 root root 4.0M Apr  2 21:55 fourmb.txt
```

[illegible][illegible]

Now let us copy the `fivemb.txt` file to the device:

As expected, no space left on device error is thrown. Next to check how many bytes are written to the device by checking the `dmesg`:

Task 3

1. List arguments of lseek function and explain their meanings

List of the arguments is:

```
loff_t fourMB_lseek(struct file *file, loff_t offset, int whence);
```

Meanings:

file: The file whose current file offset we want to change.

offset: The amount the byte offset is to be changed. The sign indicates whether the offset is to be moved forward (positive) or backward (negative).

Whence: controls the behavior and it could have three possible values:

- 1) **SEEK_SET** : the file offset should be set to offset bytes.
- 2) **SEEK_CUR** : the file offset should be set to its current location plus offset
- 3) **SEEK_END**: the file offset should be set to the size of the file (in our case is current number of bytes in the device driver) plus offset.

2. Run the test program and give screenshot of the results

There is a small change I made to the test.c:

```
printf("unable to open lcd");  
exit(EXIT_FAILURE);
```

I changed `printk` to `printf` in order to compile it to `test.o`:

```
root@yangmo-ubuntu:~/a4/task3# gcc test.c -o test.o
```

Run the test and get the results:

```
root@yangmo-ubuntu:~/a4/task3# ./test.o  
lseek = 0  
test begin!  
lseek = 4  
written = 3  
lseek = 10  
lseek = 6  
lseek = -1  
root@yangmo-ubuntu:~/a4/task3# cat /dev/lcd  
1111222111root@yangmo-ubuntu:~/a4/task3#
```

The reason why `lseek` returns 10 for "`k = lseek(lcd, 0, SEEK_END);`" after writing 3 bytes is that before writing another `lseek` function set the pointer to offset 4, which then cause the 3 bytes writing function to overwrite the old value from offset 4. Hence the total bytes are still 10 and this is why `SEEK_END` with offset 0 will return value 10.

Task 4

1. Basic Hello function and show hello message

There is one extra line needed to be added to define `SCULL_IOC_MAXNR` in addition to the sample hello code:

```
#define SCULL_IOC_MAGIC 'k'
#define SCULL_HELLO _IO(SCULL_IOC_MAGIC, 1)
#define SCULL_IOC_MAXNR 1
```

Run the test program as following and check the `dmesg`:

```
root@yangmo-ubuntu:~/a4/task4# ./test.o
test begin!
written = 3
result = 0
root@yangmo-ubuntu:~/a4/task4# dmesg
[34553.537088] Total Bytes written so far: 3
[34553.537095] hello
root@yangmo-ubuntu:~/a4/task4#
```

Hello was printed as expected.

2. Implement and test `_IOW` and `_IOR`

Added cases in `ioctl` method to handle `WRITE` and `READ`

```
case SCULL_WRITE_MESSAGE:
    max_length = MAX_DEV_MSG_LENGTH;
    msg = (char *)arg;
    tmp_dev_msg = dev_msg;
    while(*msg && max_length > 1){
        copy_from_user(tmp_dev_msg++, msg++, sizeof(char));
        max_length --;
    }
    *tmp_dev_msg = '\0';
    printk(KERN_WARNING "message written: %s\n", dev_msg);
    break;
case SCULL_READ_MESSAGE:
    msg = (char *)arg;
    tmp_dev_msg = dev_msg;
    while(*tmp_dev_msg){
        copy_to_user(msg++, tmp_dev_msg++, sizeof(char));
    }

    //terminate the input message
    put_user('\0', msg);
    break;
```

It is important that in read and write section, we need to put '\0' in the end to terminate the char array properly so when it is being printed, it could be printed properly as expected.

I also made an assumption here that the max length of the `dev_msg` is 1000:

```
//set the dev msg max length to be 1000
#define MAX_DEV_MSG_LENGTH 1000
```

I added following test cases in the `test.c`:

```
int k, i, sum;
char s[3];
char *msg = "Yang Mo A0091836X";
char msg_read[30];

memset(s, '2', sizeof(s));
printf("test begin!\n");

k = write(lcd, s, sizeof(s));
printf("written = %d\n", k);

k = ioctl(lcd, SCULL_HELLO);
printf("result = %d\n", k);

k = ioctl(lcd, SCULL_WRITE_MESSAGE, msg);
printf("result = %d\n", k);

k = ioctl(lcd, SCULL_READ_MESSAGE, msg_read);
printf("result = %d\n", k);
printf("result = %s\n", msg_read);
```

Run the test and get following result and check `dmesg` for IOW method:

```
root@yangmo-ubuntu:~/a4/task4# ./test.o
test begin!
written = 3
result = 0
result = 0
result = 0
result = Yang Mo A0091836X
root@yangmo-ubuntu:~/a4/task4# dmesg
[39404.348954] Total Bytes written so far: 3
[39404.348962] hello
[39404.348965] message written: Yang Mo A0091836X
```

3. Implement and test `_IOWR`

We need to add an `original_dev_msg` to store the old value of `dev_msg`

```
char *original_dev_msg = NULL;
```

As shown in next page, before overwriting the `dev_msg`, we need to copy its old value to `original_dev_msg`, the data of which will be copied to the user message param later before the method returns.

```

case SCULL_WRITE_READ_MESSAGE:
    strncpy(original_dev_msg, dev_msg, sizeof(char)*MAX_DEV_MSG_LENGTH);
    max_length = MAX_DEV_MSG_LENGTH;
    msg = (char *)arg;
    tmp_dev_msg = dev_msg;
    while(*msg && max_length){
        copy_from_user(tmp_dev_msg++, msg++, sizeof(char));
        max_length --;
    }
    *tmp_dev_msg = '\0';
    printk(KERN_WARNING "new dev_msg written: %s\n", dev_msg);

    //now read original value back
    msg = (char *)arg;
    tmp_dev_msg = original_dev_msg;
    while(*tmp_dev_msg){
        copy_to_user(msg++, tmp_dev_msg++, sizeof(char));
    }

    //terminate the input message
    put_user('\0', msg);

    break;

```

Then we again need to update the `test.c` to test the new IOWR method:

```

int k, i, sum;
char s[3];
char *msg = "Yang Mo A0091836X";
char msg_read[30];
char newMsg[30] = "This is a new msg";

memset(s, '2', sizeof(s));
printf("test begin!\n");

k = write(lcd, s, sizeof(s));
printf("written = %d\n", k);

k = ioctl(lcd, SCULL_HELLO);
printf("result = %d\n", k);

k = ioctl(lcd, SCULL_WRITE_MESSAGE, msg);
printf("result = %d\n", k);

k = ioctl(lcd, SCULL_READ_MESSAGE, msg_read);
printf("result = %d\n", k);
printf("result = %s\n", msg_read);

k = ioctl(lcd, SCULL_WRITE_READ_MESSAGE, newMsg);
printf("result = %d\n", k);
printf("Original Message = %s\n", newMsg);

```

Now let's run the test to display old value and check the `dmesg` for new value:

```





















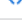





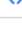




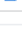
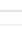
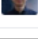























root@yangmo-ubuntu:~/a4/task4# ./test.o
test begin!
written = 3
result = 0
result = 0
result = 0
result = Yang Mo A0091836X
result = 0
Original Message = Yang Mo A0091836X
root@yangmo-ubuntu:~/a4/task4# dmesg
[45634.155537] Total Bytes written so far: 3
[45634.155547] hello
[45634.155550] message written: Yang Mo A0091836X
[45634.155556] new dev_msg written: This is a new msg

```

As can be seen, the original value was successfully returned back for display while the new value has been successfully written to `dev_msg` and logged by `printk`.

Appendix A

Github link: <https://github.com/ymoctavia/CS5250/commits/master>

Commits on Apr 4, 2017		
 fix magic number in init ymoctavia committed on GitHub 8 minutes ago	 b3fc47b	
Commits on Apr 3, 2017		
 added IOWR ymoctavia committed on GitHub 32 minutes ago	 70a5f9c	
 remove max_length > 1 check ymoctavia committed on GitHub 2 hours ago	 211dc5	
 fix typo for data termination ymoctavia committed on GitHub 2 hours ago	 5836d3d	
 terminate input data in IOW ymoctavia committed on GitHub 2 hours ago	 142ac6c	
 add ioctl Write and Read ymoctavia committed on GitHub 2 hours ago	 4080f2d	
 Add IOW and IOR declarations ymoctavia committed on GitHub 3 hours ago	 9202f4a	
 add simple ioctl function ymoctavia committed on GitHub 3 hours ago	 5f594ca	
Commits on Apr 2, 2017		
 cast offset to int32_t in lseek ymoctavia committed on GitHub a day ago	 01dcbb8	
 reset important variables in lseek function to enable correct write a... ymoctavia committed on GitHub a day ago	 d24b135	
 changed read and write logic based on f_pos param ymoctavia committed on GitHub a day ago	 c480692	
 Add lseek implementation - initial version ymoctavia committed on GitHub a day ago	 ea90436	
Commits on Apr 1, 2017		
 remove printk which was for testing purpose ymoctavia committed on GitHub 3 days ago	 86ebae8	
 move 4MB limit checker before actual data copying ymoctavia committed on GitHub 3 days ago	 5248a50	
Commits on Mar 31, 2017		
 add data length recorder and total bytes counter for both read and write ymoctavia committed on GitHub 3 days ago	 88b1e14	
 changed fourMB_write function ymoctavia committed on GitHub 3 days ago	 e67a852	
Commits on Mar 30, 2017		
 Add variables and logic to handle read and write larger files ymoctavia committed on GitHub 5 days ago	 f60e4d7	
Commits on Mar 29, 2017		
 Rename methods and variables to fourMB ymoctavia committed on GitHub 5 days ago	 4430b4c	
 Copy from Last Assignment as starting point ymoctavia committed on GitHub 5 days ago	 dbebd3b	

Appendix B

Code Link: <https://github.com/ymoctavia/CS5250/blob/master/Assignment%204/FourMBDevice.c>

```
#include <linux/module.h>
#include <linux/kernel.h>
#include <linux/init.h>
#include <linux/slab.h>
#include <linux/errno.h>
#include <linux/types.h>
#include <linux/fs.h>
#include <linux/proc_fs.h>
#include <asm/uaccess.h>
#include <linux/ioctl.h> /* needed for the _IOW etc stuff used later*/

#define SCULL_IOC_MAGIC 'k'
#define SCULL_HELLO _IO(SCULL_IOC_MAGIC, 1)
#define SCULL_WRITE_MESSAGE _IOW(SCULL_IOC_MAGIC, 2, char *)
#define SCULL_READ_MESSAGE _IOR(SCULL_IOC_MAGIC, 3, char *)
#define SCULL_WRITE_READ_MESSAGE _IOWR(SCULL_IOC_MAGIC, 4, char *)

#define SCULL_IOC_MAXNR 4

//set the dev msg max length to be 1000
#define MAX_DEV_MSG_LENGTH 1000

#define MAJOR_NUMBER 61/* forward declaration */

int fourMB_open(struct inode *inode, struct file *filep);
int fourMB_release(struct inode *inode, struct file *filep);
ssize_t fourMB_read(struct file *filep, char *buf, size_t count, loff_t *f_pos);
ssize_t fourMB_write(struct file *filep, const char *buf, size_t count, loff_t *f_pos);
loff_t fourMB_llseek(struct file *file, loff_t offset, int whence);
long fourMB_ioctl(struct file *filp, unsigned int cmd, unsigned long arg);
static void fourMB_exit(void);

/* definition of file_operation structure */
struct file_operations fourMB_fops = {
    read: fourMB_read,
    write: fourMB_write,
    open: fourMB_open,
    release: fourMB_release,
    llseek: fourMB_llseek,
    unlocked_ioctl: fourMB_ioctl,
};

char *fourMB_data = NULL;

char *data_pointer = NULL;
size_t data_length_written = 0;
size_t data_length_to_read = 0;
int bytes_written_total = 0;
int bytes_read_total = 0;
```

```

char *dev_msg = NULL;
char *original_dev_msg = NULL;

loff_t fourMB_lseek(struct file *file, loff_t offset, int whence) {
    loff_t new_position = 0;

    switch(whence) {
        case SEEK_SET :
            new_position = (int32_t)offset;
            break;
        case SEEK_CUR :
            new_position = file->f_pos + (int32_t)offset;
            break;
        case SEEK_END :
            new_position = data_length_written + (int32_t)offset;
            break;
    }

    //check boundary
    if(new_position > data_length_written){
        new_position = data_length_written;
    }

    if(new_position < 0){
        new_position = 0;
    }

    file->f_pos = new_position;

    //reset data pointer
    data_pointer = fourMB_data;

    //reset important variables
    data_length_to_read = data_length_written;
    bytes_written_total = 0;
    bytes_read_total = 0;

    return new_position;
}

long fourMB_ioctl(struct file *filp, unsigned int cmd, unsigned long arg)
{
    int err = 0;
    int retval = 0;
    char *msg;
    char *tmp_dev_msg;
    int max_length;

    /*
     * extract the type and number bitfields, and don't decode

```

```

* wrong cmds: return ENOTTY (inappropriate ioctl) before access_ok()
*/

if (_IOC_TYPE(cmd) != SCULL_IOC_MAGIC) return -ENOTTY;
if (_IOC_NR(cmd) > SCULL_IOC_MAXNR) return -ENOTTY;

/*
 * the direction is a bitmask, and VERIFY_WRITE catches R/W
 * transfers. 'Type' is user-oriented, while
 * access_ok is kernel-oriented, so the concept of "read" and
 * "write" is reversed
 */

if (_IOC_DIR(cmd) & _IOC_READ)
    err = !access_ok(VERIFY_WRITE, (void __user *)arg, _IOC_SIZE(cmd));
else if (_IOC_DIR(cmd) & _IOC_WRITE)
    err = !access_ok(VERIFY_READ, (void __user *)arg, _IOC_SIZE(cmd));

if (err) return -EFAULT;

switch(cmd) {
    case SCULL_HELLO:
        printk(KERN_WARNING "hello\n");
        break;
    case SCULL_WRITE_MESSAGE:
        max_length = MAX_DEV_MSG_LENGTH;
        msg = (char *)arg;
        tmp_dev_msg = dev_msg;
        while(*msg && max_length > 1){
            copy_from_user(tmp_dev_msg++, msg++, sizeof(char));
            max_length --;
        }
        *tmp_dev_msg = '\0';
        printk(KERN_WARNING "message written: %s\n", dev_msg);
        break;
    case SCULL_READ_MESSAGE:
        msg = (char *)arg;
        tmp_dev_msg = dev_msg;
        while(*tmp_dev_msg){
            copy_to_user(msg++, tmp_dev_msg++, sizeof(char));
        }

        //terminate the input message
        put_user('\0', msg);
        break;
    case SCULL_WRITE_READ_MESSAGE:
        strncpy(original_dev_msg, dev_msg, sizeof(char)*MAX_DEV_MSG_LENGTH);
        max_length = MAX_DEV_MSG_LENGTH;
        msg = (char *)arg;
        tmp_dev_msg = dev_msg;
        while(*msg && max_length){

```

```

        copy_from_user(tmp_dev_msg++, msg++, sizeof(char));
        max_length --;
    }
    *tmp_dev_msg = '\0';
    printk(KERN_WARNING "new dev_msg written: %s\n", dev_msg);

    //now read original value back
    msg = (char *)arg;
    tmp_dev_msg = original_dev_msg;
    while(*tmp_dev_msg){
        copy_to_user(msg++, tmp_dev_msg++, sizeof(char));
    }

    //terminate the input message
    put_user('\0', msg);

    break;
default: /* redundant, as cmd was checked against MAXNR */
    return -ENOTTY;
}

return retval;
}

int fourMB_open(struct inode *inode, struct file *filep)
{
    //re-align data pointer to the start of data section
    data_pointer = fourMB_data;

    //need to set how many bytes to read for read operation
    data_length_to_read = data_length_written;

    bytes_written_total = 0;
    bytes_read_total = 0;

    return 0; // always successful
}

int fourMB_release(struct inode *inode, struct file *filep)
{
    return 0; // always successful
}

ssize_t fourMB_read(struct file *filep, char *buf, size_t count, loff_t *f_pos)
{
    int bytes_read = 0;

    /* Check if it the end of the data section */
    if (data_length_to_read == 0){

```

```

        printk("Total Bytes read: %d ", bytes_read_total);
        return 0;
    }

    data_pointer = *f_pos + data_pointer;
    data_length_to_read -= *f_pos;

    while (count && *data_pointer && data_length_to_read) {

        copy_to_user(buf++, data_pointer++, sizeof(char));

        count--;
        bytes_read++;
        bytes_read_total ++;
        data_length_to_read --;
    }

    return bytes_read;
}

ssize_t fourMB_write(struct file *filep, const char *buf, size_t count, loff_t *f_pos)
{
    int bytes_written = 0;

    //if this function is called first time during current write operation
    //simple reset the written data length
    if(data_pointer == fourMB_data){
        //printk("Offset: %d ", *f_pos);
        data_pointer += *f_pos;
    }

    while (count && *buf) {
        //detect if have written 4MB data
        if(data_length_written >= 1024*1024*4*sizeof(char)){
            break;
        }

        copy_from_user(data_pointer++, buf++, sizeof(char));

        count--;
        bytes_written++;
        bytes_written_total++;

        if((bytes_written_total + *f_pos) > data_length_written){
            data_length_written ++;
        }
    }

    printk("Total Bytes written so far: %d ", bytes_written_total);
}

```

```

        //check if data more than 4MB left to be written
        if(count > 0)
        {
            printk(KERN_ALERT "No space left on device\n");

            /*Return Linux System Error<28>: No space left on device */
            return -ENOSPC;
        }

        return bytes_written;
    }
}

static int fourMB_init(void)
{
    int result;
    // register the device
    result = register_chrdev(MAJOR_NUMBER, "fourMB", &fourMB_fops);
    if (result < 0) {
        return result;
    }
    // allocate 4 MB of memory for storage
    // kmalloc is just like malloc, the second parameter is
    // the type of memory to be allocated.
    // To release the memory allocated by kmalloc, use kfree.
    fourMB_data = kmalloc(1024*1024*4*sizeof(char), GFP_KERNEL);

    dev_msg = kmalloc(MAX_DEV_MSG_LENGTH*sizeof(char), GFP_KERNEL);
    original_dev_msg = kmalloc(MAX_DEV_MSG_LENGTH*sizeof(char), GFP_KERNEL);

    if (!fourMB_data) {
        fourMB_exit();
        // cannot allocate memory
        // return no memory error, negative signify a failure
        return -ENOMEM;
    }

    // initialize the value to be X
    *fourMB_data = 'X';
    data_length_written++;

    printk(KERN_ALERT "This is a fourMB device module\n");
    return 0;
}

static void fourMB_exit(void)
{

```

```

{
    // if the pointer is pointing to something
    if (fourMB_data) {
        // free the memory and assign the pointer to NULL
        kfree(fourMB_data);
        fourMB_data = NULL;
    }

    if (dev_msg) {
        kfree(dev_msg);
        dev_msg = NULL;
    }
    if (original_dev_msg) {
        kfree(original_dev_msg);
        original_dev_msg = NULL;
    }

    // unregister the device
    unregister_chrdev(MAJOR_NUMBER, "fourMB");
    printk(KERN_ALERT "FourMB device module is unloaded\n");
}

MODULE_LICENSE("GPL");
module_init(fourMB_init);
module_exit(fourMB_exit);

```